


# XLNet: Generalized Autoregressive Pretraining for Language Understanding



모두의 연구소 풀잎스쿨

NLP bootcamp 5<sup>th</sup> 박성찬

2019. 11. 30(sat)

# I. Introduction: XLNet

- # A SOTA pretrained model for language understanding
- # Generalized AR pretraining model
  - Permutation language modeling objectives
  - two-stream attention mechanism
- # 여러 NLP task에서 original BERT 대비 우수한 성능
- # 논문 이외의 여러 블로그 참고
  - <https://novdov.github.io/machinelearning/nlp/2019/07/13/XLNet-%EB%A6%AC%EB%B7%B0/>
  - <https://ratsgo.github.io/natural%20language%20processing/2019/09/11/xlnet/>
  - <https://blog.pingpong.us/xlnet-review/>
  - <https://ai-information.blogspot.com/2019/07/nl-041-xlnet-generalized-autoregressive.html>

# I. Introduction: Pretraining


## # How it works

- Pretrain a model on unlabeled data based on language modeling
- Finetune the model or use the model for feature extraction on downstream tasks

Word2vec(Mikolov et al.), GloVe(Pennington et),  
Semi-supervised sequence learning(Dai & Le),  
ELMo(Peters, et al.), CoVe(McCann et al.),  
GPT(Radford et al.), BERT(Devlin et al.)

## 2.1 Background: Two Objectives for Pretraining

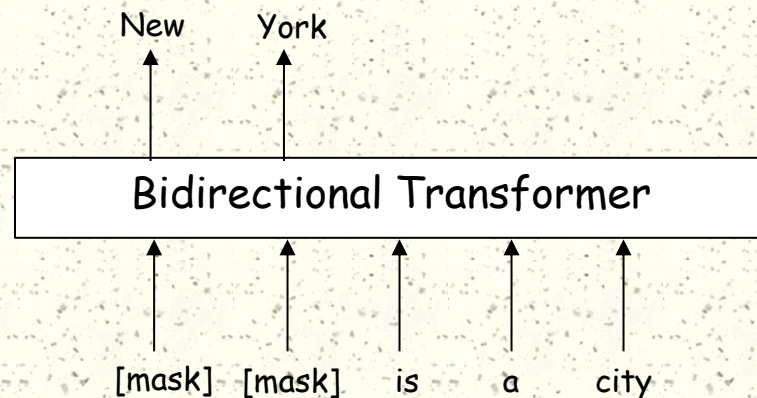
Auto-regressive(AR)  
language modeling



$$\log p(\mathbf{x}) = \sum_{t=1}^T \log p(x_t | \mathbf{x}_{<t})$$

- 양방향 context 학습 불가

(Denoising) Auto-encoding(AE)



$$\log p(\mathbf{x}^- | \mathbf{x}^+) = \sum_{t=1}^T \text{mask}_t \log p(x_t | \mathbf{x}^+)$$

- Masked token들 사이는 독립이라는 가정  
- Finetuning 과정에는 mask token 없음



## 2.2 Objective: Permutation Language Modeling(PLM)

- # Sample a factorization order  $z$
- # Determine the attention masks based on the order
- # Optimize a standard language modeling objective

$$\mathbb{E}_{z \sim Z_T} \left[ \sum_{t=1}^T \log p(x_{z_t} | \mathbf{x}_{z < t}) \right]$$

$$\begin{aligned} \text{Ex.) } p(a, b) &= p(a)p(b/a) \\ &= p(b)p(a/b) \end{aligned}$$

## 2.2 PLM example

# Sentence: New York is a city

# Factorization order: **New York is a city**

$p(\text{New York is a city}) = p(\text{New}) * p(\text{York} | \text{New}) * p(\text{is} | \text{New York}) * p(\text{a} | \text{New York is}) * p(\text{city} | \text{New York is a})$

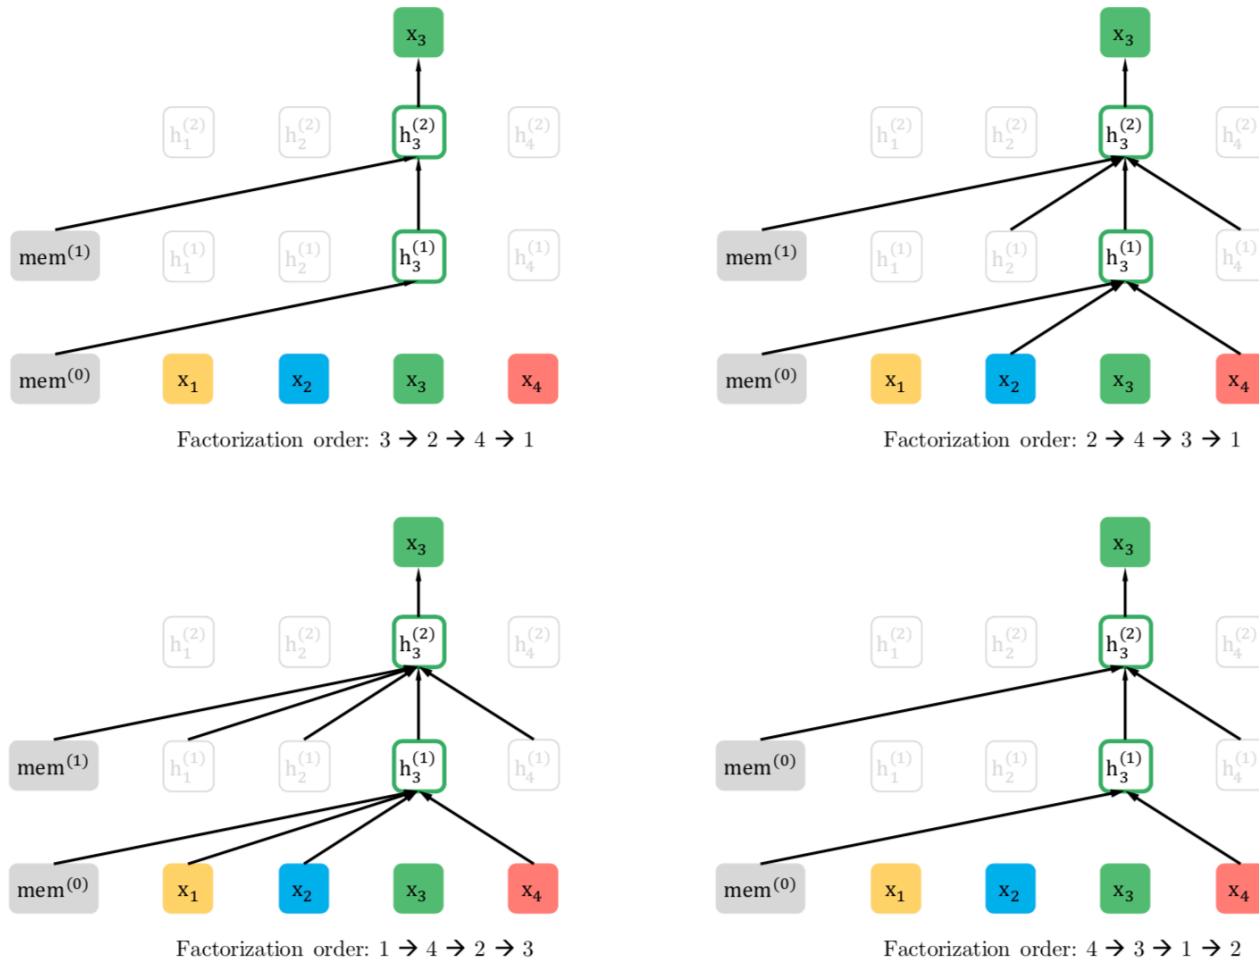
# Factorization order: **city a is New York**

$p(\text{New York is a city}) = p(\text{city}) * p(\text{a} | \text{city}) * p(\text{is} | \text{city, a}) * p(\text{New} | \text{city a is}) * p(\text{York} | \text{city a is New})$

# Remarks on Permutation

- Sequence order is not shuffled: 기존 sequence order는 유지하고 positional encoding만 바꾸어 transformer의 attention mask에 적용
- Attention masks are changed to reflect factorization order

## 2.2 PLM illustration



*Fig1: Illustration of the permutation language modeling(PLM) objective for predicting  $x_3$ , given the same input sequence  $x$  but with different factorization orders*

## 2.2 PLM objectives

- ✦ Not only the benefits of AR models, but capture bidirectional contexts

$\mathbb{Z}_t$  is set of all possible permutations of length- $t$  index seq.  $[1, 2, \dots, T]$

$z_t$  is  $t$ -th element and  $z < t$  is the first  $t-1$  elements of a permutation  $z \in \mathbb{Z}_t$

$$\max_{\theta} \mathbb{E}_{z \sim \mathcal{Z}_T} \left[ \sum_{t=1}^T \log p_{\theta}(x_{z_t} | x_{z < t}) \right] \quad (3)$$

For text sequence  $x$ , sample a factorization order  $z$  and decompose  $p_{\theta}(x)$

Same model parameter  $\theta$  is shared across all factorization orders during training, it could possibly capture the bidirectional context



## 2.3 Two-Stream Self-Attention for Target-Aware Representation: Reparameterization

### ✦ Standard Parameterization

$$p_{\theta}(X_{Z_t} = x | \mathbf{x}_{z < t}) = \frac{e(x)^T h_{\theta}(\mathbf{x}_{z < t})}{\sum_{x'} e(x')^T h_{\theta}(\mathbf{x}_{z < t})}$$

$h$  does not contain the position of the target

For input sequence  $[x_1, x_2, x_3, x_4]$ , its permutation  $Z_T = [[1, 2, 3, 4], [1, 3, 2, 4], \dots, [4, 3, 2, 1]]$

$[2, 3, 1, 4]: p(x_1 | x_2, x_3)$  → Ambiguity!

$[2, 3, 4, 1]: p(x_4 | x_2, x_3)$

### ✦ Solutions: condition the distribution on the position

$$p_{\theta}(X_{Z_t} = x | \mathbf{x}_{z < t}) = \frac{e(x)^T g_{\theta}(\mathbf{x}_{z < t}, \mathbf{z}_t)}{\sum_{x'} e(x')^T g_{\theta}(\mathbf{x}_{z < t}, \mathbf{z}_t)}$$

→ But how to compute  $g_{\theta}(\mathbf{x}_{z < t}, \mathbf{z}_t)$  effectively?

## 2.3 Two-Stream Self-Attention for Target-Aware Representation: Two-Stream Self-Attention

- # If predict token  $X_{z_t}$ ,  $g_{\theta}(x_{z< t}, z_t)$ ,
  - it should use the position  $z_t$  and not the context  $X_{z_t} \Rightarrow g_{\theta}(x_{z< t}, z_t)$ , abbreviated  $g_{z_t}$
  - $g_{z_t}^{(m)} \leftarrow \text{Attention}(Q = g_{z_t}^{(m-1)}, KV = h_{z< t}^{(m-1)}; \theta)$
  - Query representation: access the contextual information  $x_{z< t}$  and position  $z_t$  but cannot see  $X_{z_t}$
- # Otherwise,
  - it should encode content  $X_{z_t} \Rightarrow h_{\theta}(x_{z< t})$ , abbreviated  $h_{z_t}$
  - $h_{z_t}^{(m)} \leftarrow \text{Attention}(Q = h_{z_t}^{(m-1)}, KV = h_{z\leq t}^{(m-1)}; \theta)$
  - Content representation: encodes both the context and  $X_{z_t}$  itself

## 2.3 Two-Stream Self-Attention for Target-Aware Representation: Illustration1

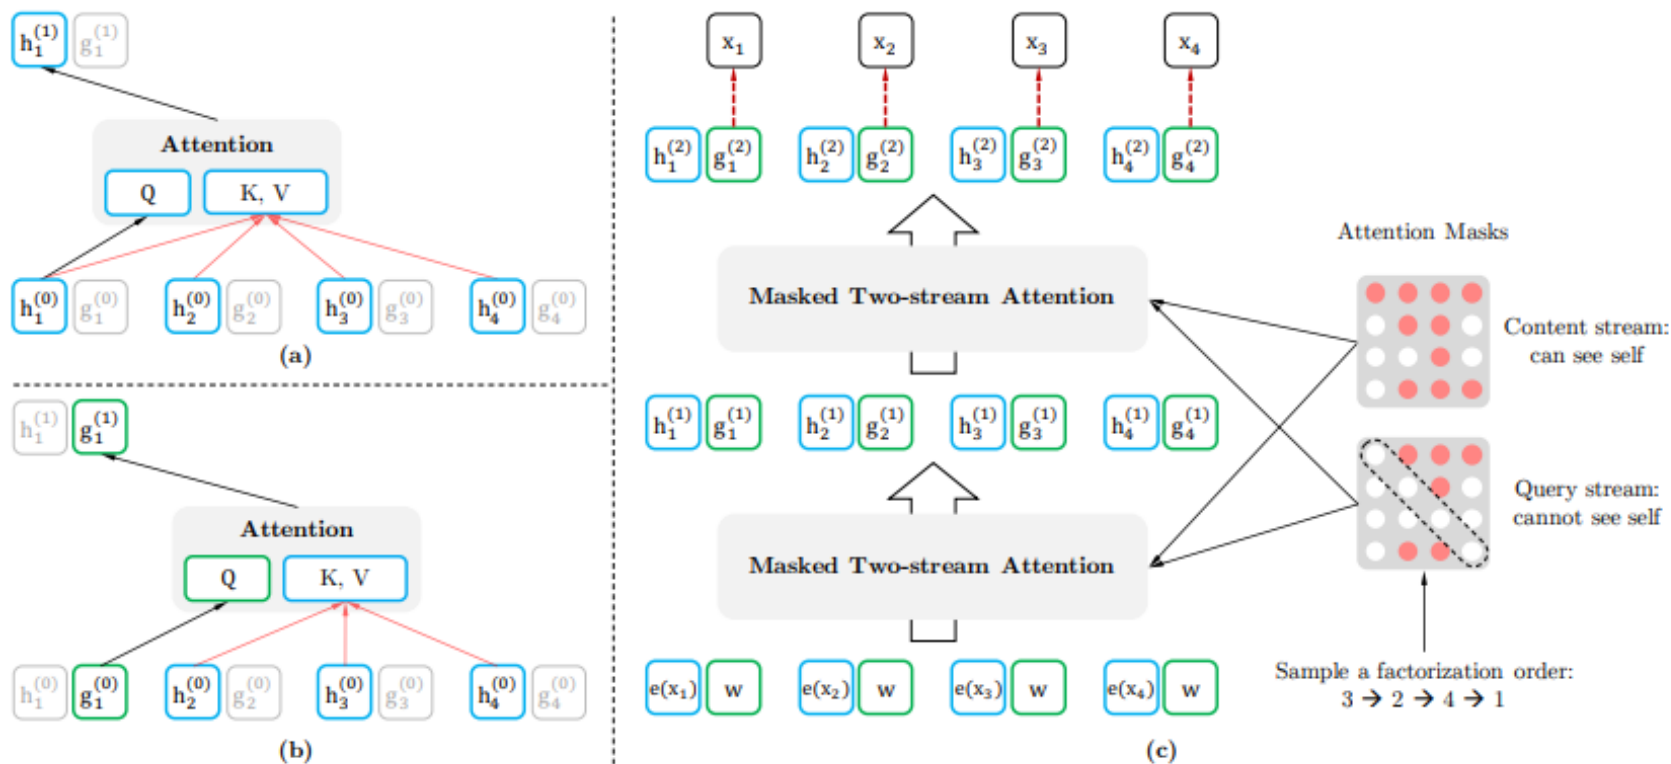




Figure 2: (a): Content stream attention, which is the same as the standard self-attention. (b): Query stream attention, which does not have access information about the content  $x_{z_t}$ . (c): Overview of the permutation language modeling training with two-stream attention.


## 2.3 Two-Stream Self-Attention for Target-




Joint View of the Content Stream  
(Factorization order:  $3 \rightarrow 2 \rightarrow 4 \rightarrow 1$ )




Position-3 View



Position-2 View

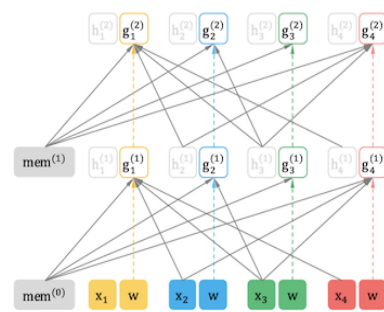


Position-4 View





Position-1 View


Split View of the Content Stream  
(Factorization order:  $3 \rightarrow 2 \rightarrow 4 \rightarrow 1$ )




Joint View of the Query Stream  
(Factorization order:  $3 \rightarrow 2 \rightarrow 4 \rightarrow 1$ )




Position-3 View



Position-2 View



Position-4 View



Position-1 View

Split View of the Query Stream  
(Factorization order:  $3 \rightarrow 2 \rightarrow 4 \rightarrow 1$ )



## 2.3 Two-Stream Self-Attention for Target-Aware Representation: Partial Prediction

✦ Several benefits, but optimization problem due to permutations

- Only predict the last tokens in a factorization order

- Split  $z$  into a non-target subsequence  $Z_{\leq c}$  and a target subsequence  $Z_{>c}$  ( $c$ : cutting point)

- $$\begin{aligned} \max_{\theta} \mathbb{E}_{z \sim Z_T} [\log p_{\theta}(x_{z_{>c}} | x_{z_{\leq c}})] = \\ \max_{\theta} \mathbb{E}_{z \sim Z_T} [\sum_{t=c+1}^T \log p_{\theta}(x_{z_t} | x_{z_{<t}})] \end{aligned}$$

- $x_{z_{>c}}$  is chosen as a target because it possesses the longest context in the sequence given the current factorization order  $Z$

- $3 \rightarrow 2 \rightarrow 4 \rightarrow 1$  order

- $p(x_3)p(x_2|x_3)p(x_4|x_2, x_3)p(x_1|x_3, x_2, x_4) \rightarrow p(x_4|x_2, x_3)p(x_1|x_3, x_2, x_4)$

## 2.4 Incorporating Ideas from Transformer-XL

### # Relative positional encoding scheme

- 여러 segment에 대해 recurrent 모델링하는 경우 absolute positional encoding은 더이상 사용 불가

### # Segment recurrent mechanism

- Long sentence  $A$  divides  $\tilde{x} = s_{1:T}, x = s_{T+1:2T}$  and let  $\tilde{z} = \text{perm}[1, \dots, T], z = \text{perm}[T+1, \dots, 2T]$
- Process the first segment based on  $\tilde{z}$  then cache the obtained content representations  $\tilde{h}^{(m)}$  each layer  $m$
- $h_{z_t}^{(m)} \leftarrow \text{Attention}(Q = h_{z_t}^{(m-1)}, KV = [\tilde{h}^{(m-1)}, h^{(m-1)}_{z \leq t}]; \theta)$

## 2.5 Modeling Multiple Segments

### # XLNet Pre-training

- 임의로 두 **segment**를 뽑아 하나로 **concatenation**하여 PLM
- 동일한 **context**에 속해 있는 **memory**는 재사용
- Bert와 유사하게 **[A, SEP, B, SEP, CLS]** 형태(A, B는 **segment**)의 입력

### # Relative segment encodings

- BERT: word embedding + abs segment embedding
- XLNet: relative segment embedding
  - **Sequence**의 위치  $i, j$ 에 대하여 어떠한 **segment**에서 왔는지는 생각 안 하고 **s vector**를 사용하여 동일한 **segment**인지만 고려
  - Finetuning 단계에서 두 개가 넘는 **segment**입력을 받을 수 있음
  - Inductive bias of relative encoding의 일반화?

## 2.6.1 Comparing with BERT: concepts

# Sentence: New York is a city

# BERT objective(auto-encoding)

$$\mathcal{J}_{BERT} = \log p(\text{New}|\text{is a city}) + \log p(\text{York}|\text{is a city})$$

- New와 York은 독립 사건

# XLNet objective(auto-regressive)

$$\mathcal{J}_{XLNet} = \log p(\text{New}|\text{is a city}) + \log p(\text{York}|\text{New}, \text{is a city})$$

$$\mathcal{J}_{XLNet} = \log p(\text{New}|\text{York}, \text{is a city}) + \log p(\text{York}|\text{is a city})$$

- New 와 York사이의 dependency 모델링

- Joint Probability를 factorization하여 bidirectional context 모델링

- 더 많은 dependency 학습 가능("denser" effective training signals)



## 2.6.1 Comparing with BERT: formal expressions

- ✦ Given sequence  $x=[x_1, x_2, \dots, x_T]$ , target-context pair of interest  $\mathcal{L} = \{(x, \mathcal{U})\}$

$\mathcal{L} = \{(x = \text{York} | \mathcal{U} = \{\text{New}\}), (x = \text{York} | \mathcal{U} = \{\text{city}\}), (x = \text{york}, \mathcal{U} = \{\text{New}, \text{city}\}), \dots\}$

- ✦  $\mathcal{T}$  is given target tokens,  $\mathcal{N} = x \setminus \mathcal{T}$  is non-target tokens

- Both BERT & XLNet maximize  $\log p(\mathcal{T} | \mathcal{N})$

$$\mathcal{J}_{BERT} = \sum_{x \in \mathcal{T}} \log p(x | \mathcal{N}); \quad \mathcal{J}_{XLNet} = \sum_{x \in \mathcal{T}} \log p(x | \mathcal{N} \cup \mathcal{T}_{<x})$$

- Both has multiple loss terms  $\log p(x | \mathcal{V}_x)$ ,  $\mathcal{V}_x = \mathcal{N}$  or  $\mathcal{N} \cup \mathcal{T}_{<x}$

- ✦ Two cases by definition

- If  $\mathcal{U} \in \mathcal{N}$ , the dependency  $(x, \mathcal{U})$  covered by both
- If  $\mathcal{U} \in \mathcal{N} \cup \mathcal{T}_{<x}$  and  $\mathcal{N} \cup \mathcal{T}_{<x} \neq \emptyset$ , the dependency  $(x, \mathcal{U})$  covered only by XLNet

## 2.6.2 Comparing with language modeling

### # Conventional AR language modeling

- Context: "Thom Yorke is the singer of Radiohead"
- Question: "Who is the singer of Radiohead"
- Answer: "Thom York"

- AR은 뒤에 나오는 문맥과의 dependency 학습이 어려움

### # Formal expression

- Consider a context-target pair  $(x, u)$
- If  $u \cap \mathcal{T}_{<x} \neq \emptyset$ , where  $\mathcal{T}_{<x}$  denotes tokens prior to  $x$  in the original sentence, AR cannot able to cover the dependency

- XLNet은 모든 dependency 학습 가능

# 3.1 Pretraining and Implementation

## # Dataset

- BookCorpus + English Wikipedia = 13GB
- Giga5=16GB
- Clue Web2012-B+Common Crawl dataset = 19GB, 78GB

## # Tokenization

- Using sentence piece, it sums total 32.89B subword tokens

## # Pretraining

- Sequence length = 512, memory length = 384
- XLNet large= 512 TPU v3 chips for 500k steps with Adam optimizer, batch size 2048, about 2.5 days

## 3.2-3.3 Race & SQuAD Dataset

RACE	Accuracy	Middle	High
GPT [25]	59.0	62.9	57.4
BERT [22]	72.0	76.6	70.1
BERT+OCN* [28]	73.5	78.4	71.5
BERT+DCMN* [39]	74.1	79.5	71.8
XLNet	<b>81.75</b>	<b>85.45</b>	<b>80.21</b>

Table 1: Comparison with state-of-the-art results on the test set of RACE, a reading comprehension task. \* indicates using ensembles. “Middle” and “High” in RACE are two subsets representing middle and high school difficulty levels. All BERT and XLNet results are obtained with a 24-layer architecture with similar model sizes (aka BERT-Large). Our single model outperforms the best ensemble by 7.6 points in accuracy.

SQuAD1.1	EM	F1	SQuAD2.0	EM	F1
<i>Dev set results without data augmentation</i>					
BERT [10]	84.1	90.9	BERT† [10]	78.98	81.77
XLNet	<b>88.95</b>	<b>94.52</b>	XLNet	<b>86.12</b>	<b>88.79</b>
<i>Test set results on leaderboard, with data augmentation (as of June 19, 2019)</i>					
Human [27]	82.30	91.22	BERT+N-Gram+Self-Training [10]	85.15	87.72
ATB	86.94	92.64	SG-Net	85.23	87.93
BERT* [10]	87.43	93.16	BERT+DAE+AoA	85.88	88.62
XLNet	<b>89.90</b>	<b>95.08</b>	XLNet	<b>86.35</b>	<b>89.13</b>

Table 2: A single model XLNet outperforms human and the best ensemble by 7.6 EM and 2.5 EM on SQuAD1.1. \* means ensembles, † marks our runs with the official code.



# 3.4-3.5 Text Classification & GLUE dataset

Model	IMDB	Yelp-2	Yelp-5	DBpedia	AG	Amazon-2	Amazon-5
CNN [14]	-	2.90	32.39	0.84	6.57	3.79	36.24
DPCNN [14]	-	2.64	30.58	0.88	6.87	3.32	34.81
Mixed VAT [30, 20]	4.32	-	-	0.70	4.95	-	-
ULMFIT [13]	4.6	2.16	29.98	0.80	5.01	-	-
BERT [35]	4.51	1.89	29.32	0.64	-	2.63	34.17
XLNet	<b>3.79</b>	<b>1.55</b>	<b>27.80</b>	<b>0.62</b>	<b>4.49</b>	<b>2.40</b>	<b>32.26</b>

Table 3: Comparison with state-of-the-art error rates on the test sets of several text classification datasets. All BERT and XLNet results are obtained with a 24-layer architecture with similar model sizes (aka BERT-Large).

Model	MNLI	QNLI	QQP	RTE	SST-2	MRPC	CoLA	STS-B	WNLI
<i>Single-task single models on dev</i>									
BERT [2]	86.6/-	92.3	91.3	70.4	93.2	88.0	60.6	90.0	-
XLNet	<b>89.8/-</b>	<b>93.9</b>	<b>91.8</b>	<b>83.8</b>	<b>95.6</b>	<b>89.2</b>	<b>63.6</b>	<b>91.8</b>	-
<i>Single-task single models on test</i>									
BERT [10]	86.7/85.9	91.1	89.3	70.1	94.9	89.3	60.5	87.6	65.1
<i>Multi-task ensembles on test (from leaderboard as of June 19, 2019)</i>									
Snorkel* [29]	87.6/87.2	93.9	89.9	80.9	96.2	91.5	63.8	90.1	65.1
ALICE*	88.2/87.9	95.7	<b>90.7</b>	83.5	95.2	92.6	<b>68.6</b>	91.1	80.8
MT-DNN* [18]	87.9/87.4	96.0	89.9	<b>86.3</b>	96.5	92.7	68.4	91.1	89.0
XLNet*	<b>90.2/89.7<sup>†</sup></b>	<b>98.6<sup>†</sup></b>	90.3 <sup>†</sup>	<b>86.3</b>	<b>96.8<sup>†</sup></b>	<b>93.0</b>	67.8	<b>91.6</b>	<b>90.4</b>

Table 4: Results on GLUE. \* indicates using ensembles, and <sup>†</sup> denotes single-task results in a multi-task row. All results are based on a 24-layer architecture with similar model sizes (aka BERT-Large). See the upper-most rows for direct comparison with BERT and the lower-most rows for comparison with state-of-the-art results on the public leaderboard.

## 3.6-3.7 ClueWeb09-B dataset & Ablation study

Model	NDCG@20	ERR@20
DRMM [12]	24.3	13.8
KNRM [8]	26.9	14.9
Conv [8]	28.7	18.1
BERT <sup>†</sup>	30.53	18.67
XLNet	<b>31.10</b>	<b>20.28</b>

Table 5: Comparison with state-of-the-art results on the test set of ClueWeb09-B, a document ranking task. <sup>†</sup> indicates our implementations.

#	Model	RACE	SQuAD2.0		MNLI m/mm	SST-2
			F1	EM		
1	BERT-Base	64.3	76.30	73.66	84.34/84.65	92.78
2	DAE + Transformer-XL	65.03	79.56	76.80	84.88/84.45	92.60
3	XLNet-Base ( $K = 7$ )	66.05	<b>81.33</b>	<b>78.46</b>	<b>85.84/85.43</b>	92.66
4	XLNet-Base ( $K = 6$ )	66.66	80.98	78.18	85.63/85.12	<b>93.35</b>
5	- memory	65.55	80.15	77.27	85.32/85.05	92.78
6	- span-based pred	65.95	80.61	77.91	85.49/85.02	93.12
7	- bidirectional data	66.34	80.65	77.87	85.31/84.99	92.66
8	+ next-sent pred	<b>66.76</b>	79.83	76.94	85.32/85.09	92.89

Table 6: Ablation study. The results of BERT on RACE are taken from [39]. We run BERT on the other datasets using the official implementation and the same hyperparameter search space as XLNet.  $K$  is a hyperparameter to control the optimization difficulty (see Section 2.3). All models are pretrained on the same data.

## 4. Conclusion

---

XLNet is a generalized AR pretraining method

Permutation language modeling

Combine the advantages of AR and AE methods

XLNet architecture

Transformer XL & two-stream attention mechanism

XLNet achieves SOTA results from various tasks